

Image Classification with Pre-trained Models (Keras Applications)

Keras provides many pre-trained models trained on the ImageNet dataset (1.4M images, 1000 classes).

These models can be used for:\

- Image classification
- Feature extraction
- Fine-tuning for custom tasks.

Keras Applications are state-of-the-art, pre-trained CNNs (trained on ImageNet, 1k classes).

They're great for:\

- Prediction (zero-shot image classification)
- Feature extraction (as backbones for your tasks)
- Fine-tuning (transfer learning on your dataset)

Import Libraries

- time → measure how fast predictions are\
- numpy → for handling images as numbers\
- tensorflow → the deep learning library\
- image → load and preprocess images\
- Pre-trained models & their helpers → ResNet, Xception, EfficientNet, NASNet

```
In [7]: import time
import numpy as np
import tensorflow as tf
from tensorflow.keras.preprocessing import image

# Model families (for preprocess & decode)
from tensorflow.keras.applications import resnet_v2, xception, efficientnet, nas

# Individual models
from tensorflow.keras.applications.resnet_v2 import ResNet50V2
from tensorflow.keras.applications.xception import Xception
from tensorflow.keras.applications.efficientnet import EfficientNetB3, Efficient
from tensorflow.keras.applications.nasnet import NASNetLarge

print("TensorFlow version:", tf.__version__)
```

TensorFlow version: 2.20.0

Load & Preprocess an Image

Goal: Prepare our image for the model.

Why: Models expect numbers in a specific format (size, batch, normalization)

```
In [10]: def load_and_preprocess(img_path, target_size, preprocess_fn):
        """
        1. Load image from disk
        2. Resize to model's expected size
        3. Convert to a numerical array
        4. Add a batch dimension (model expects 1 image at a time)
        5. Apply model-specific preprocessing
        """
        img = image.load_img(img_path, target_size=target_size)
        x = image.img_to_array(img)
        x = np.expand_dims(x, axis=0)
        x = preprocess_fn(x)
        return x
```

Set Image Path

```
In [12]: IMG_PATH = r"C:\Ds & AI ( my work)\Deep_Learning\Image_Classification_with_Pre-t
import os
print("Using image:", IMG_PATH, "| Exists:", os.path.exists(IMG_PATH))
```

Using image: C:\Ds & AI (my work)\Deep_Learning\Image_Classification_with_Pre-trained_Models_(Keras)\input-2.jpeg | Exists: True

```
In [14]: #Load image
img = image.load_img(IMG_PATH)
img
```

Out[14]:



Run One Model at a Time

1.ResNet50V2

- Why ResNet50V2?\
 - Reliable, fast, good accuracy.\
 - Input size: 224×224 pixels.\
- What happens here:\
 1. Image resized to 224×224 → model expects this.
 2. predict() → gives probabilities for 1000 ImageNet classes.

3. decode_predictions() → converts numbers to readable labels, e.g., "jeep", "car", etc

```
In [18]: # Load and preprocess image
x = load_and_preprocess(IMG_PATH, target_size=(224,224),
    preprocess_fn=resnet_v2.preprocess_input)

# Load pre-trained model
model = ResNet50V2(weights="imagenet")

# Predict
start = time.time()
preds = model.predict(x)
end = time.time()

# Decode predictions (top-5)
from tensorflow.keras.applications.resnet_v2 import decode_predictions
top5 = decode_predictions(preds, top=5)[0]

print(f"Inference time: {end-start:.3f} sec")
print("Top-5 Predictions:")
for i, (imagenet_id, label, score) in enumerate(top5, 1):
    print(f"{i}. {label:25s} ({score:.2f})")
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet50v2_weights_tf_dim_ordering_tf_kernels.h5

102869336/102869336 ————— 39s 0us/step

1/1 ————— 3s 3s/step

Downloading data from https://storage.googleapis.com/download.tensorflow.org/data/imagenet_class_index.json

35363/35363 ————— 0s 5us/step

Inference time: 3.568 sec

Top-5 Predictions:

1. Eskimo_dog	(0.33)
2. Siberian_husky	(0.21)
3. Pomeranian	(0.07)
4. malamute	(0.05)
5. Samoyed	(0.04)

2. Xception

- Why Xception?
 - Uses depthwise separable convolutions → efficient and accurate.
 - Input size: 299×299 pixels.

```
In [22]: x = load_and_preprocess(IMG_PATH, target_size=(299,299), preprocess_fn=xception.
    preprocess_input)
model = Xception(weights="imagenet")
preds = model.predict(x)
top5 = xception.decode_predictions(preds, top=5)[0]

print("Top-5 Predictions (Xception):")
for i, (imagenet_id, label, score) in enumerate(top5, 1):
    print(f"{i}. {label:25s} ({score:.2f})")
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/xception/xception_weights_tf_dim_ordering_tf_kernels.h5

91884032/91884032 ————— 27s 0us/step

1/1 ————— 2s 2s/step

Top-5 Predictions (Xception):

- | | |
|---------------------|--------|
| 1. Pomeranian | (0.45) |
| 2. Chihuahua | (0.27) |
| 3. papillon | (0.04) |
| 4. Japanese_spaniel | (0.01) |
| 5. toy_terrier | (0.01) |

3. NASNetLarge

- Why NASNetLarge?
 - Very accurate, designed via neural architecture search.
 - Input: 331×331 pixels.

```
In [23]: x = load_and_preprocess(IMG_PATH, target_size=(331,331), preprocess_fn=nasnet.
         preprocess_input)
         model = NASNetLarge(weights="imagenet")
         preds = model.predict(x)
         top5 = nasnet.decode_predictions(preds, top=5)[0]

         print("Top-5 Predictions (NASNetLarge):")
         for i, (imagenet_id, label, score) in enumerate(top5, 1):
             print(f"{i}. {label:25s} ({score:.2f})")
```

Downloading data from <https://storage.googleapis.com/tensorflow/keras-applications/nasnet/NASNet-large.h5>

359748576/359748576 ————— 144s 0us/step

1/1 ————— 11s 11s/step

Top-5 Predictions (NASNetLarge):

- | | |
|---------------------|--------|
| 1. Pomeranian | (0.76) |
| 2. Chihuahua | (0.08) |
| 3. papillon | (0.01) |
| 4. keeshond | (0.00) |
| 5. Japanese_spaniel | (0.00) |

Key Takeaways (Non-Tech Version)

1. ResNet50V2 → fast, reliable baseline.
2. Xception → efficient, accurate with medium input size.
3. NASNetLarge → very accurate but slow.

Tip: For beginners or real-time apps → use ResNet50V2 or EfficientNetB3. For research/accuracy comparison → try EfficientNetB7 or NASNetLarge.

Conclusion & Summary

- Pre-trained Models are deep learning models already trained on large datasets like ImageNet.

- They can be used for image classification, feature extraction, or transfer learning without training from scratch.
- Transformers and CNNs like EfficientNet, ResNet, Xception, and NASNet are popular pre-trained models.
- **How to use them:**
 1. Load the model with pre-trained weights.\
 2. Preprocess your image (resize, normalize) according to the model's requirements.\
 3. Run inference to predict classes.\
 4. Optionally, fine-tune the model for custom datasets.\
- Real-time applications & examples:–
 - mobile apps: Recognizing objects or animals in photos.
 - Self-driving cars: Detecting pedestrians, vehicles, traffic signs.
 - Healthcare: Identifying diseases from X-rays or scans.
 - Retail: Visual search, product recognition, and recommendation.
- **Model choice tip:** - Lightweight models like ResNet50V2 or EfficientNetB3 → fast predictions for real time apps.
 - Heavier models like NASNetLarge or EfficientNetB7 → higher accuracy, used in research or offline processing.
- Pre-trained models save time, computation, and improve accuracy while enabling practical real-world applications

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []: